

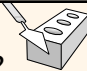
The Relational Model

Chapter 3

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

1

1



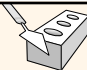
Why Study the Relational Model?

- ❖ Most widely used model.
 - Vendors: IBM, Microsoft, Oracle, etc.
 - Open Source: e.g., MySQL, and PostgreSQL
- ❖ Former competitors:
 - The Network Model: Now legacy
 - The Hierarchical Model: Now legacy
 - Object-oriented Model
 - ObjectStore, O2, Poet, ..
- ❖ Ingested into: *object-relational model*
 - Supported by all industrial-strength systems

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

2

2



Relational Database: Definitions

- ❖ **Relational database**: a set of **relations**
- ❖ **Relation**: made up of 2 parts:
 - **Instance** : a **table**, with rows and columns.
#Rows = **cardinality**, #fields = **degree / arity**.
 - **Schema** : specifies name of relation, plus name and type of each column.
 - E.G. Students(sid: string, name: string, login: string, age: integer, gpa: real).
- ❖ Can think of a relation as a **set** of rows or **tuples** (i.e., all rows are distinct).

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

3

3

Example Instance of Students Relation

| sid | name | login | age | gpa |
|-------|-------|------------|-----|-----|
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@eecs | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |

- ❖ Cardinality = 3, degree = 5, all rows distinct
- ❖ Do all columns in a relation instance have to be distinct?

4

Relational Query Languages

- ❖ A major strength of the relational model: supports simple, powerful *querying* of data.
- ❖ Queries can be written intuitively, and the DBMS is responsible for efficient evaluation.
 - The key: precise semantics for relational queries.
 - Allows the optimizer to extensively re-order operations, and still ensure that the answer does not change.

5

The SQL Query Language

- ❖ First developed by IBM (system R) in the 1970s
- ❖ Need for a standard since it is used by many vendors
- ❖ Standards:
 - SQL-86: By 1986, ANSI and ISO standard groups officially adopted the standard "Database Language SQL" language definition.
 - New versions of the standard were published in 1989, 1992, 1996, 1999, 2003, 2006, 2008, 2011, and, 2016.

6

sql is declarative language and will translated into relational algebra and query evaluation pipes by machine

Relational algebra is a procedural query language

Relational Calculus is a Declarative language (tuple and domain relational calculus)

The SQL Query Language

- ❖ To find all 18 year old students, we can write:

```
SELECT * FROM Students S
WHERE S.age=18
```

Student S is a tuple variable
where clause select predicate (with one variable)
join (two variables)

- To find just names and logins, replace the first line:

```
SELECT S.name, S.login
```

Querying Multiple Relations

- ❖ What does the following query compute?

```
SELECT S.name, E.cid
FROM Students S, Enrolled E
WHERE S.sid=E.sid AND E.grade="A"
```

get rid of S. E. will still work because these attributes are unique

Given the following instances of Enrolled and Students:

| sid | name | login | age | gpa |
|-------|-------|------------|-----|-----|
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@ecs | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |

| sid | cid | grade |
|-------|-------------|-------|
| 53831 | Carnatic101 | C |
| 53831 | Reggae203 | B |
| 53650 | Topology112 | A |
| 53666 | History105 | B |

we get:

| S.name | E.cid |
|--------|-------------|
| Smith | Topology112 |

self join:

Parent child table (foreign key)

Child | Parent

select Child and Grand Parent

Select p2.c, p1.p
from parents p1, p2
where p1.c = p2.p

Creating Relations in SQL

- ❖ Creates the Students relation. Observe that the type (domain) of each field is specified, and enforced by the DBMS whenever tuples are added or modified.

```
CREATE TABLE Students
(sid: CHAR(20),
 name: CHAR(20),
 login: CHAR(10),
 age: INTEGER,
 gpa: REAL)
```

- ❖ As another example, the Enrolled table holds information about courses that students take.

```
CREATE TABLE Enrolled
(sid: CHAR(20),
 cid: CHAR(20),
 grade: CHAR(2))
```

Destroying and Altering Relations

DROP TABLE Students

- ❖ Destroys the relation Students. The schema information *and* the tuples are deleted.

ALTER TABLE Students

ADD COLUMN firstYear: integer

- ❖ The schema of Students is altered by adding a new field; every tuple in the current instance is extended with a *null* value in the new field.

10

Adding and Deleting Tuples

- ❖ Can insert a single tuple using:

```
INSERT INTO Students (sid, name, login, age, gpa)
VALUES (53688, 'Smith', 'smith@ee', 18, 3.2)
```

- ❖ Can delete all tuples satisfying some condition (e.g., name = Smith):

```
DELETE
FROM Students S
WHERE S.name = 'Smith'
```

bulk delete (delete all)

bulk update (update all)

☛ Powerful variants of these commands are available; more later!

11

Integrity Constraints (ICs)

- ❖ **IC**: condition that must be true for *any* instance of the database; e.g., domain constraints.

- ICs are specified when schema is defined.
- ICs are checked when relations are modified.

- ❖ A *legal* instance of a relation is one that satisfies all specified ICs.

- DBMS should not allow illegal instances.

- ❖ If the DBMS checks ICs, stored data is more faithful to real-world meaning.

- Avoids data entry errors, too!

12

primary key has to be unique, email should have @, e.g

data type + Constraints (NOT NULL / UNIQUE / PRIMARY KEY / FOREIGN KEY / CHECK / DEFAULT)

Primary Key Constraints

- ❖ A set of fields is a key for a relation if : candidate keys: keys that can be chose from
 1. No two distinct tuples can have same values in all key fields, and
 2. This is not true for any subset of the key.
 - Part 2 false? A superkey.
 - If there's >1 key for a relation, one of the keys is chosen (by DBA) to be the primary key. the key that we chose
- ❖ E.g., *sid* is a key for Students. (What about *name*?) The set {*sid*, *gpa*} is a superkey.

13

Primary and Candidate Keys in SQL

- ❖ Possibly many candidate keys (specified using UNIQUE), one of which is chosen as the primary key.
- ❖ "For a given student and course, there is a single grade." vs. "Students can take only one course, and receive a single grade for that course; further, no two students in a course receive the same grade."

```
CREATE TABLE Enrolled
(sid CHAR(20)
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid) )
```
- ❖ Used carelessly, an IC can prevent the storage of database instances that arise in practice!

```
CREATE TABLE Enrolled
(sid CHAR(20)
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid),
UNIQUE (cid, grade) )
```

14

Foreign Keys, Referential Integrity

- ❖ Foreign key : Set of fields in one relation that is used to 'refer' to a tuple in another relation. (Must correspond to primary key of the second relation.) Like a 'logical pointer'.
- ❖ E.g. *sid* is a foreign key referring to Students:
 - Enrolled(*sid*: string, *cid*: string, *grade*: string)
 - If all foreign key constraints are enforced, referential integrity is achieved, i.e., no dangling references.
 - Can you name a data model w/o referential integrity?
 - Links in HTML!

15

Foreign Keys in SQL

- ❖ Only students listed in the Students relation should be allowed to enroll for courses.

```
CREATE TABLE Enrolled
(sid CHAR(20), cid CHAR(20), grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid) REFERENCES Students )
```

Enrolled

| sid | cid | grade |
|-------|-------------|-------|
| 53666 | Carnatic101 | C |
| 53666 | Reggae203 | B |
| 53650 | Topology112 | A |
| 53666 | History105 | B |

Students

| sid | name | login | age | gpa |
|-------|-------|------------|-----|-----|
| 53666 | Jones | jones@cs | 18 | 3.4 |
| 53688 | Smith | smith@eecs | 18 | 3.2 |
| 53650 | Smith | smith@math | 19 | 3.8 |

16

Enforcing Referential Integrity

- ❖ Consider Students and Enrolled; *sid* in Enrolled is a foreign key that references Students.
- ❖ What should be done if an Enrolled tuple with a non-existent student id is inserted? (*Reject it!*)
- ❖ What should be done if a Students tuple is deleted?
 - Also delete all Enrolled tuples that refer to it.
 - Disallow deletion of a Students tuple that is referred to.
 - Set *sid* in Enrolled tuples that refer to it to a *default sid*.
 - (In SQL, also: Set *sid* in Enrolled tuples that refer to it to a special value *null*, denoting 'unknown' or 'inapplicable'.)
- ❖ Similar if primary key of Students tuple is updated.

17

Referential Integrity in SQL

- ❖ SQL/92 and SQL:1999 support all 4 options on deletes and updates.
 - Default is **NO ACTION** (*delete/update is rejected*)
 - **CASCADE** (also delete all tuples that refer to deleted tuple)
 - **SET NULL / SET DEFAULT** (sets foreign key value of referencing tuple)

```
CREATE TABLE Enrolled
(sid CHAR(20),
cid CHAR(20),
grade CHAR(2),
PRIMARY KEY (sid,cid),
FOREIGN KEY (sid)
REFERENCES Students
ON DELETE CASCADE
ON UPDATE SET DEFAULT )
```

18

what if students dropped purdue and what happened to enrolled table

- 1 on delete cascade (delete tuples from other tables with same foreign keys)
- 2 on delete set null
- 3 on delete restrict (reject is it has references)
- 4 on delete set default (set the value to default value)


Where do ICs Come From?

- ❖ ICs are based upon the semantics of the real-world enterprise that is being described in the database relations.
- ❖ We can check a database instance to see if an IC is violated, but we can **NEVER** infer that an IC is true by looking at an instance.
 - An IC is a statement about *all possible* instances!
 - From example, we know *name* is not a key, but the assertion that *sid* is a key is given to us.
- ❖ Key and foreign key ICs are the most common; more general ICs supported too.

19

Logical DB Design: Mapping to the Relational Model

- ❖ Map entities to tables:



```
CREATE TABLE Employees
(ssn CHAR(11),
name CHAR(20),
lot INTEGER,
PRIMARY KEY (ssn))
```

20

Relationship Sets to Tables

- ❖ In translating a relationship set to a relation, attributes of the relation must include:
 - Keys for each participating entity set (as foreign keys).
 - This set of attributes forms a *superkey* for the relation.
 - All descriptive attributes.

```
CREATE TABLE Works_In(
ssn CHAR(11),
did INTEGER,
since DATE,
PRIMARY KEY (ssn, did),
FOREIGN KEY (ssn)
REFERENCES Employees,
FOREIGN KEY (did)
REFERENCES Departments)
```

21

Key Constraints

binary relation

❖ Each dept has at most one manager, according to the key constraint on Manages.

Translation to relational model?

1-to-1 1-to Many Many-to-1 Many-to-Many

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke 22

22

Translating ER Diagrams with Key Constraints

❖ Map relationship to a table:

- Note that **did** is the key now!
- Separate tables for Employees and Departments.

❖ Since each department has a unique manager, we could instead combine Manages and Departments.

```
CREATE TABLE Manages(
  ssn CHAR(11),
  did INTEGER,
  since DATE,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn) REFERENCES Employees,
  FOREIGN KEY (did) REFERENCES Departments)

CREATE TABLE Dept_Mgr(
  did INTEGER,
  dname CHAR(20),
  budget REAL,
  ssn CHAR(11),
  since DATE,
  PRIMARY KEY (did),
  FOREIGN KEY (ssn) REFERENCES Employees)
```

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke 23

23


Participation Constraints

❖ Does every department have a manager?

- If so, this is a participation constraint: the participation of Departments in Manages is said to be total (vs. partial).
- Every **did** value in Departments table must appear in a row of the Manages table (with a non-null **ssn** value!)

arrow: every dept has only one manager

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke 24



24

Participation Constraints in SQL

- ❖ We can capture participation constraints involving one entity set in a binary relationship, but little else (without resorting to CHECK constraints).

```
CREATE TABLE Dept_Mgr(  
  did INTEGER,  
  dname CHAR(20),  
  budget REAL,  
  ssn CHAR(11) NOT NULL,  
  since DATE,  
  PRIMARY KEY (did),  
  FOREIGN KEY (ssn) REFERENCES Employees,  
  ON DELETE NO ACTION)
```

25

Views

- ❖ A view is just a relation, but we store a *definition*, rather than a set of tuples.

```
CREATE VIEW YoungActiveStudents (name, grade)  
AS SELECT S.name, E.grade  
FROM Students S, Enrolled E  
WHERE S.sid = E.sid and S.age < 21
```

- ❖ Views can be dropped using the **DROP VIEW** command.
 - How to handle **DROP TABLE** if there's a view on the table?
 - DROP TABLE command has options to let the user specify this.

26

Views and Security

- ❖ Views can be used to present necessary information (or a summary), while hiding details in underlying relation(s).
 - Given YoungStudents, but not Students or Enrolled, we can find students s who have are enrolled, but not the *cid*'s of the courses they are enrolled in.

27

Relational Model: Summary



- ❖ A tabular representation of data.
- ❖ Simple and intuitive, currently the most widely used.
- ❖ Integrity constraints can be specified by the DBA, based on application semantics. DBMS checks for violations.
 - Two important ICs: primary and foreign keys
 - In addition, we *always* have domain constraints.
- ❖ Powerful and natural query languages exist.
- ❖ Rules to translate a real world data application scenario into the relational model
