

## Evaluation of Relational Operations

### Chapter 14, Part A (Joins)

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

1

---

---

---


---

---

---

---

1



## Relational Operations

- ❖ We will consider how to implement:
  - Selection ( $\sigma$ ) Selects a subset of rows from relation.
  - Projection ( $\pi$ ) Deletes unwanted columns from relation.
  - Join ( $\bowtie$ ) Allows us to combine two relations.
  - Set-difference ( $-$ ) Tuples in reln. 1, but not in reln. 2.
  - Union ( $\cup$ ) Tuples in reln. 1 and in reln. 2.
  - Aggregation (SUM, MIN, etc.) and GROUP BY
- ❖ Since each op returns a relation, ops can be *composed*! After we cover the operations, we will discuss how to *optimize* queries formed by composing them.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

2

---

---

---


---

---

---

---

2



## Schema for Examples

Sailors (sid: integer, sname: string, rating: integer, age: real)  
 Reserves (sid: integer, bid: integer, day: dates, rname: string)

- ❖ Similar to old schema; *rname* added for variations.
- ❖ Reserves:
  - Each tuple is 40 bytes long, 100 tuples per page, 1000 pages.
- ❖ Sailors:
  - Each tuple is 50 bytes long, 80 tuples per page, 500 pages.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

3

---

---

---

---

---

---

---

3

## Equality Joins With One Join Column

```
SELECT *  
FROM Reserves R1, Sailors S1  
WHERE R1.sid=S1.sid
```

- ❖ In algebra:  $R \bowtie S$ . Common! Must be carefully optimized.  $R \times S$  is large; so,  $R \times S$  followed by a selection is inefficient.
- ❖ Assume:  $M$  pages in  $R$ ,  $p_R$  tuples per page,  $N$  tuples in  $S$ ,  $p_S$  tuples per page.
  - In our examples,  $R$  is Reserves and  $S$  is Sailors.
- ❖ We will consider more complex join conditions later.
- ❖ **Cost metric:** # of I/Os. We will ignore output costs.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

4

4

---

---

---

---

---

---

---

## Simple Nested Loops Join

```
foreach tuple r in R do  
  foreach tuple s in S do  
    if  $r_i == s_j$  then add  $\langle r, s \rangle$  to result
```

- ❖ For each tuple in the *outer* relation  $R$ , we scan the entire *inner* relation  $S$ .
  - **Cost:**  $M + p_R * M * N = 1000 + 100 * 1000 * 500$  I/Os.
- ❖ Page-oriented Nested Loops join: For each *page* of  $R$ , get each *page* of  $S$ , and write out matching pairs of tuples  $\langle r, s \rangle$ , where  $r$  is in  $R$ -page and  $S$  is in  $S$ -page.
  - **Cost:**  $M + M * N = 1000 + 1000 * 500$
  - If smaller relation ( $S$ ) is outer, cost =  $500 + 500 * 1000$

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

5

5

---

---

---

---

---

---

---

## Index Nested Loops Join

```
foreach tuple r in R do  
  foreach tuple s in S where  $r_i == s_j$  do  
    add  $\langle r, s \rangle$  to result
```

- ❖ If there is an index on the join column of one relation (say  $S$ ), can make it the inner and exploit the index.
  - **Cost:**  $M + (M * p_R) * \text{cost of finding matching } S \text{ tuples}$
- ❖ For each  $R$  tuple, cost of probing  $S$  index is about 1.2 for hash index, 2-4 for B+ tree. Cost of then finding  $S$  tuples (assuming Alt. (2) or (3) for data entries) depends on clustering. (additional I/O to get tuple from tuple id stored in leaf level)
  - **Clustered index:** 1 I/O (typical), **unclustered:** upto 1 I/O per matching  $S$  tuple.

6

6

Assuming 80% occupancy for hash index, then the size of hash index is 1.2\*the size of the table.

Thus the cost of probing is about 1.2

Hash index here is cheaper because equal join

---

---

---

---

---

---

---

## Examples of Index Nested Loops



### ❖ Hash-index (Alt. 2) on *sid* of Sailors (as inner):

- Scan Reserves: 1000 page I/Os, 100\*1000 tuples.
- For each Reserves tuple: 1.2 I/Os to get data entry in index, plus 1 I/O to get (the exactly one) matching Sailors tuple. Total: 220,000 I/Os.

### ❖ Hash-index (Alt. 2) on *sid* of Reserves (as inner):

- Scan Sailors: 500 page I/Os, 80\*500 tuples.
- For each Sailors tuple: 1.2 I/Os to find index page with data entries, plus cost of retrieving matching Reserves tuples. Assuming uniform distribution, 2.5 reservations per sailor (100,000 / 40,000). Cost of retrieving them is 1 or 2.5 I/Os depending on whether the index is clustered.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

7

7

---

---

---

---

---

---

---


---

## Block Nested Loops Join



### ❖ Use one page as an input buffer for scanning the inner S, one page as the output buffer, and use all remaining pages to hold "block" of outer R.

- For each matching tuple *r* in R-block, *s* in S-page, add *<r, s>* to result. Then read next R-block, scan S, etc.



Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

8

8

store the block as a hash table  
scan the whole inner table once per block  
and for each tuple in the page S  
we probe the hash table and store the matching pair to output  
buffer  
once a page S is exhausted, do the same with next page

---

---

---

---

---

---

---

---

## Examples of Block Nested Loops



### ❖ Cost: Scan of outer + #outer blocks \* scan of inner

- #outer blocks =  $\lceil \# \text{ of pages of outer} / \text{blocksize} \rceil$

### ❖ With Reserves (R) as outer, and 100 pages of R: a block of size 100

- Cost of scanning R is 1000 I/Os; a total of 10 blocks.
- Per block of R, we scan Sailors (S); 10\*500 I/Os.
- If space for just 90 pages of R, we would scan S 12 times.

### ❖ With 100-page block of Sailors as outer:

- Cost of scanning S is 500 I/Os; a total of 5 blocks.
- Per block of S, we scan Reserves; 5\*1000 I/Os.

### ❖ With sequential reads considered, analysis changes: may be best to divide buffers evenly between R and S.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

9

9

---

---

---

---

---

---

---

---

## Sort-Merge Join ( $R \bowtie S$ )



- ❖ Sort R and S on the join column, then scan them to do a "merge" (on join col.), and output result tuples.
  - Advance scan of R until current R-tuple  $\geq$  current S tuple, then advance scan of S until current S-tuple  $\geq$  current R tuple; do this until current R tuple = current S tuple.
  - At this point, all R tuples with same value in  $R_i$  (current R group) and all S tuples with same value in  $S_j$  (current S group) **match**; output  $\langle r, s \rangle$  for all pairs of such tuples.
  - Then resume scanning R and S.
- ❖ R is scanned once; each S group is scanned once per matching R tuple. (Multiple scans of an S group are likely to find needed pages in buffer.)

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

10

10

assume both tables are pre-sorted

have a pointer for each table

if R.sid is smaller than S.sid, advance R.sid

if S.sid is smaller than R.sid,

advance S.sid until find a match

## Example of Sort-Merge Join



sid	sname	rating	age	sid	bid	day	rname
22	dustin	7	45.0	28	103	12/4/96	guppy
28	yuppy	9	35.0	28	103	11/3/96	yuppy
31	lubber	8	55.5	31	101	10/10/96	dustin
44	guppy	5	35.0	31	102	10/12/96	lubber
58	rusty	10	35.0	31	101	10/11/96	lubber
				58	103	11/12/96	dustin

- ❖ Cost:  $M \log M + N \log N + (M+N)$ 
  - The cost of scanning,  $M+N$ , could be  $M*N$  (very unlikely!)
- ❖ With 35, 100 or 300 buffer pages, both Reserves and Sailors can be sorted in 2 passes; total join cost: 7500.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

11

11

cost to sort 1 table + cost to sort another table + cost of merge  
(because we need to scan both tables one tuple at a time)

## Refinement of Sort-Merge Join



- ❖ We can combine the merging phases in the *sorting* of R and S with the merging required for the join.
  - With  $B > \sqrt{L}$ , where  $L$  is the size of the larger relation, using the sorting refinement that produces runs of length  $2B$  in Pass 0, #runs of each relation is  $< B/2$ .
  - Allocate 1 page per run of each relation, and "merge" while checking the join condition.
  - Cost: read+write each relation in Pass 0 + read each relation in (only) merging pass (+ writing of result tuples).
  - In example, cost goes down from 7500 to 4500 I/Os.
- ❖ In practice, cost of sort-merge join, like the cost of external sorting, is **linear**.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke


12

12


**hash-join: use a hash partition to smaller the size of relations again and again until partition can fit in main memory**

## Hash-Join

- ❖ Partition both relations using hash fn  $h$ : R tuples in partition  $i$  will only match S tuples in partition  $i$ .



- ❖ Read in a partition of R, hash it using  $h_2$  ( $\neq h$ ). Scan matching partition of S, search for matches.



Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

13

13

---

---

---

---

---

---

---

---

## Observations on Hash-Join

- ❖ #partitions  $k < B-1$  (why?), and  $B-2 > \text{size of largest partition}$  to be held in memory. Assuming uniformly sized partitions, and maximizing  $k$ , we get:
  - $k = B-1$ , and  $M/(B-1) < B-2$ , i.e.,  $B$  must be  $> \sqrt{M}$
- ❖ If we build an in-memory hash table to speed up the matching of tuples, a little more memory is needed.
- ❖ If the hash function does not partition uniformly, one or more R partitions may not fit in memory. Can apply hash-join technique recursively to do the join of this R-partition with corresponding S-partition.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

14

14

---

---

---

---

---

---

---

---

## Cost of Hash-Join

- ❖ In partitioning phase, read+write both relns;  $2(M+N)$ . In matching phase, read both relns;  $M+N$  I/Os.
- ❖ In our running example, this is a total of 4500 I/Os.
- ❖ Sort-Merge Join vs. Hash Join:
  - Given a minimum amount of memory (*what is this, for each?*) both have a cost of  $3(M+N)$  I/Os. Hash Join superior on this count if relation sizes differ greatly. Also, Hash Join shown to be highly parallelizable.
  - Sort-Merge less sensitive to data skew; result is sorted.

Database Management Systems 3ed, R. Ramakrishnan and J. Gehrke

15

15

---

---

---

---

---

---

---

---

## General Join Conditions



- ❖ Equalities over several attributes (e.g.,  $R.sid=S.sid$  AND  $R.rname=S.rname$ ):
  - For Index Nested Loop Join (index has to be built on inner table), build index on  $\langle sid, rname \rangle$  (if S is inner); or use existing indexes on  $sid$  or  $rname$ .
  - For Sort-Merge and Hash Join, sort/partition on combination of the two join columns.
- ❖ Inequality conditions (e.g.,  $R.rname < S.rname$ ):
  - For Index NL, need (clustered!) B+ tree index.
    - Range probes on inner; # matches likely to be much higher than for equality joins.
  - Hash Join, Sort Merge Join not applicable.
  - Block NL quite likely to be the best join method here.

16

16

Find the leaf node which rname is equal, then to find the names greater than it, we just get all nodes on the right (has to be clustered, otherwise tuple ids are all over the place)

