

Relational Query Processing

Walid Aref

Spring 2020

1

1

How Queries Get Processed

- Start from an SQL query:

2

2

How Queries Get Processed

- Start from an SQL query:
 - Select * from A, B, C Where
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"

User Interface
(Command Line)

3

3

How Queries Get Processed

- Start from an SQL query:
 - Select * from A, B, C Where
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"
- Transform this query to its equivalent relational algebra expression:

User Interface
(Command Line)
Query Parser and
Query Rewrite

4

4

How Queries Get Processed

- Start from an SQL query:
 - Select * from A, B, C Where
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"
- Transform this query to an equivalent relational algebra expression:
 - $\sigma_{A.a0 = "123"}(A \bowtie_{A.a1 = B.b1} B \bowtie_{B.b2 = C.c2} C)$

User Interface
(Command Line)
Query Parser and
Query Rewrite

5

5

How Queries Get Processed

- Start from an SQL query:
 - Select * from A, B, C Where
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"
- Transform this query to an equivalent relational algebra expression:
 - $\sigma_{A.a0 = "123"}(A \bowtie_{A.a1 = B.b1} B \bowtie_{B.b2 = C.c2} C)$
- Optimize this expression:

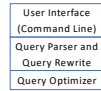
User Interface
(Command Line)
Query Parser and
Query Rewrite
Query Optimizer

6

6

How Queries Get Processed

- Start from an SQL query:
 - Select * from A, B, C Where
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"
- Transform this query to an equivalent relational algebra expression:
 - $\sigma_{A.a0 = "123"}(A \bowtie_{A.a1 = B.b1} B \bowtie_{B.b2 = C.c2} C)$
- Optimize this expression
 - $(\sigma_{A.a0 = "123"}(A) \bowtie_{A.a1 = B.b1} B) \bowtie_{B.b2 = C.c2} C$

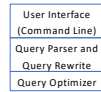


7

7

How Queries Get Processed

- Start from an SQL query:
 - Select * from A, B, C Where
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"
- Transform this query to an equivalent relational algebra expression:
 - $\sigma_{A.a0 = "123"}(A \bowtie_{A.a1 = B.b1} B \bowtie_{B.b2 = C.c2} C)$
- Optimize this expression:
 - $(\sigma_{A.a0 = "123"}(A) \bowtie_{A.a1 = B.b1} B) \bowtie_{B.b2 = C.c2} C$
- Generate an equivalent query evaluation plan:

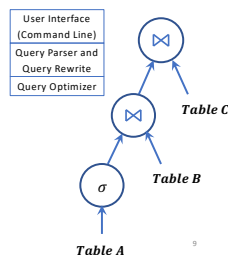


8

8

How Queries Get Processed

- Start from an SQL query:
 - Select * from A, B, C Where
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"
- Transform this query to an equivalent relational algebra expression:
 - $\sigma_{A.a0 = "123"}(A \bowtie_{A.a1 = B.b1} B \bowtie_{B.b2 = C.c2} C)$
- Optimize this expression:
 - $(\sigma_{A.a0 = "123"}(A) \bowtie_{A.a1 = B.b1} B) \bowtie_{B.b2 = C.c2} C$
- Generate an equivalent query evaluation plan:



9

9

How Queries Get Processed

- Start from an SQL query:
 - Select * from A, B, C Where
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"
- Transform this query to an equivalent relational algebra expression:
 - $\sigma_{A.a0 = "123"}(A \bowtie_{A.a1 = B.b1} B \bowtie_{B.b2 = C.c2} C)$
- Question:
 - Is the above transformation always possible?

User Interface
(Command Line)
Query Parser and
Query Rewrite



?

10

10

Conceptual Evaluation of SQL Queries into Relational Algebra Expressions

- Given an SQL query:
 - Select attr₁, attr₂, ..., attr_n
 - From Table₁, Table₂, ..., Table_m
 - Where Pred₁ and Pred₂ and ... and Pred_k
- The equivalent Relational Algebra expression:
 - $\pi_{attr_1, attr_2, \dots, attr_n} (\sigma_{Pred_1 \wedge Pred_2 \wedge \dots \wedge Pred_k} (Table_1 \times Table_2 \times \dots \times Table_m))$

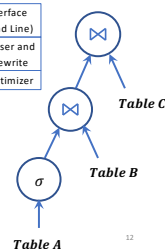
11

11

How Queries Get Processed

- Start from an SQL query:
 - Select * from A, B, C Where
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"
- Transform this query to an equivalent relational algebra expression:
 - $\sigma_{A.a0 = "123"}(A \bowtie_{A.a1 = B.b1} B \bowtie_{B.b2 = C.c2} C)$
- Optimize this expression:
 - $(\sigma_{A.a0 = "123"}(A) \bowtie_{A.a1 = B.b1} B) \bowtie_{B.b2 = C.c2} C$
- Generate an equivalent query evaluation plan:

User Interface
(Command Line)
Query Parser and
Query Rewrite
Query Optimizer



12

12

How Queries Get Processed

- Start from an SQL query:
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"

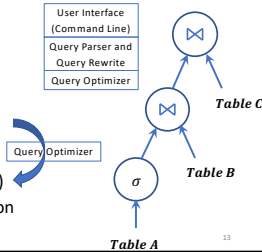
- Transform this query to an equivalent relational algebra expression:

$$\sigma_{A.a0 = "123"}(A \bowtie_{A.a1 = B.b1} B \bowtie_{B.b2 = C.c2} C)$$

- Optimize this expression:

$$((\sigma_{A.a0 = "123"}(A) \bowtie_{A.a1 = B.b1} B) \bowtie_{B.b2 = C.c2} C)$$

- Generate an equivalent query evaluation plan:



13

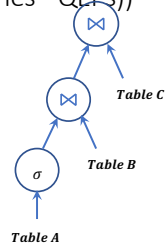
Query Evaluation Trees (or Query Evaluation Plans (Or Query Evaluation Pipelines - QEPs))

- Select * from A, B, C Where
 - A.a1 = B.b1 and B.b2 = C.c2 and
 - A.a0 = "123"

- Query Parser -> Query Rewrite -> Query Optimizer -> Query Evaluation Pipeline - Query Processor

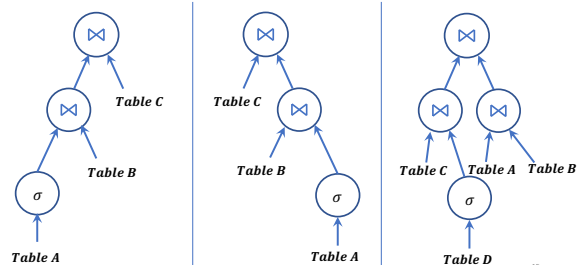
- Questions:

- How does a query evaluation pipeline work?
- Why is it called "pipeline"?
- What are the various shapes of query evaluation pipelines?
- Does the shape of a QEP matter for query processing purposes?

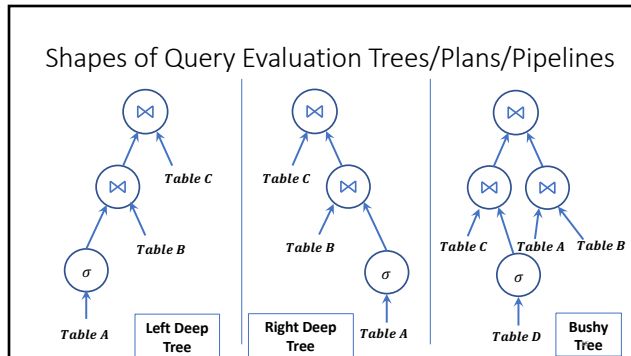


14

Shapes of Query Evaluation Trees/Plans/Pipelines



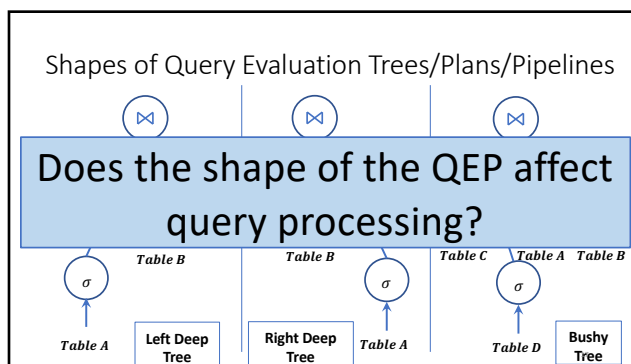
15



16

assume right-deep tree has a temporary table and temporary table has p tuples. Then the cost of IO is

$$c * p(c \text{ joins with temp table}) + p(\text{save temporary table to disk}) + ab(\text{iterate all } a, b \text{ tuples to construct temporary table})$$



17

The Iterator Model to Evaluate Relational Algebra Operators

- The iterator model:
 - Process one tuple at a time
 - Follow the lazy evaluation scheme
 - Is a pull-based model
 - Does not produce a tuple until it is asked
 - In contrast to "push-based model"
 - Produce tuples without being asked
- Each relational algebra operator is implemented using the iterator model

18

How to Code an Operator in the Iterator Model?

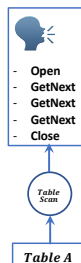
- Each operator has the following interface methods:
 - Open():**
 - Open the operator
 - Initialize the operator and its internal state variables
 - GetNext():**
 - Start from the current state of the operator
 - Execute the operator from its current state until it is able to retrieve the next tuple
 - Report the next tuple
 - Advance and store the current state for the next execution
 - If no tuples are available for return (e.g., reached end-of-table), report done/null/end-of-table
 - Close():**
 - Close the operator
 - Other auxiliary operators:
 - HasNext(), Reset() or ReOpen(), among others

19

19

Physical Operators and their Realizations as Iterators

- The Table Scan Operator:
 - Open():**
 - Open the table
 - Set the cursor at the beginning of the table
 - GetNext():**
 - Get the next tuple from the table
 - Advance the cursor to the next tuple
 - If no tuple exists, then return null or end-of-table
 - Else return the tuple
 - Close():**
 - Close the underlying table
 - Reset():**
 - Reset the cursor to point to the beginning of the table

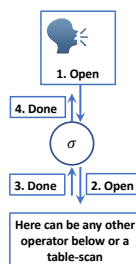


20

20

Physical Operators and their Realizations as Iterators: The Select Operator

- The Select Operator. $\sigma_P(\dots)$:
 - Upon construction at compile time, Predicate P of the select is passed to the operator to be embedded in it (or passed as a Hint. More on this later).
 - Open():**
 - Issue an Open() operator to the operator under it (could be a table scan operator or any other operator underneath)



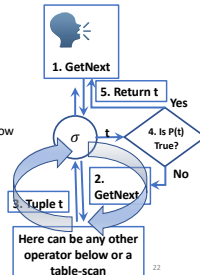
21

21

Physical Operators and their Realizations as Iterators: The Select Operator

The Select Operator:

- **getNext():**
 - Try to getNext() from the operator below the select - $\rightarrow t$
 - Check P on t
 - If true, return t and exit
 - Else repeat the above until a tuple is found or operator below reports Done.
- **Close():**
 - Call Close() for the operator below the select
 - Return done

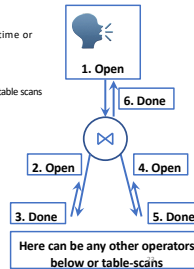


22

22

Physical Operators and their Realizations as Iterators: The Join Operator

- The Join Operator:
- The join predicate will be plugged into the inside of the operator at compile time or will be passed as a hint (more on this issue later).
- State = left_tuple and right_tuple
- **Open():**
 - Issue an Open() operator to both the left and right operators under the join (could be table scans or any other operators underneath)
 - Set left_tuple and right_tuple to be null
- **getNext():**
 - If left_tuple = null, then Get the next tuple from the left operator underneath
 - While left operator did not reach end_of_table do
 - If right_tuple = null, then Get the next tuple from the right operator underneath
 - While join_predicate is not true and not end_of_table of right operator do
 - right_tuple ← right_operator.getNext()
 - If join_predicate is true then return left_tuple join right_tuple
 - Else
 - Reset right operator
 - left_tuple ← left_operator.getNext()
- **Close():**
 - Call Close() for both the left and right operators below the join
 - Return done



23

Putting the Operators Together to Construct a Query Evaluation Pipeline

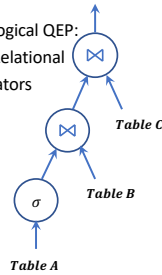
- Connect the operators together to form a QEP.
- Each operator does NOT know the operator underneath it
 - And it does not need to know.
- Each operator follows the iterator model protocol (open, getNext, close).
- Thus, operators can communicate successfully and move tuples around to answer the query.

24

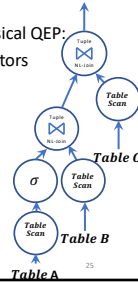
24

Transforming a Logical QEP to a Physical QEP

- An Example Logical QEP:
- Uses Logical Relational Algebra operators



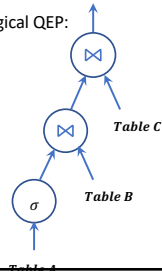
- Corresponding Physical QEP:
- Uses Physical operators



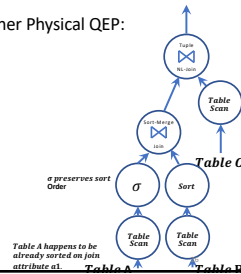
25

Transforming a Logical QEP to a Physical QEP: Using Sort-Merge Joins + Sort Operators

- An Example Logical QEP:



- Another Physical QEP:



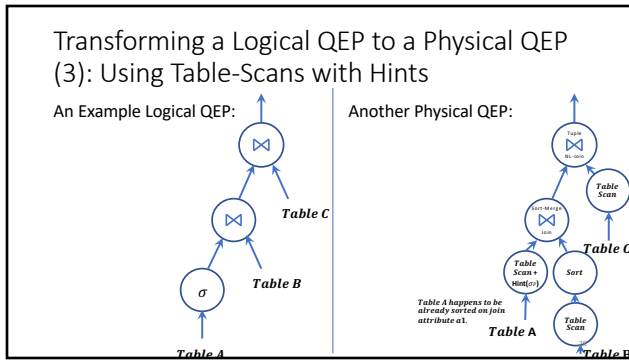
26

Physical Operators

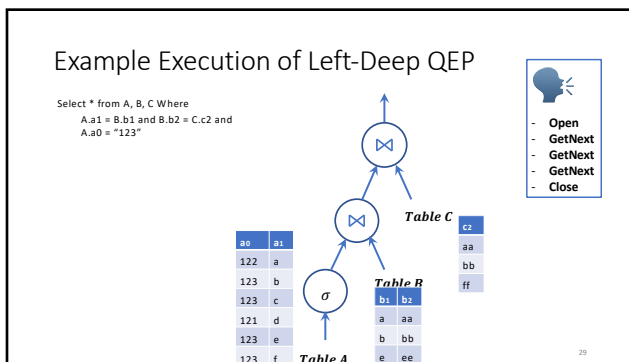
- Table-Scan
- Table-Scan with select predicate as a hint
- Index-Scan (Useful for index-only plans)
- Index-Scan with select predicate as a hint
- Temp (Create a temporary table out of the incoming tuples from the lower-level operator)
- Create-index (Create index on the fly)
- Table-Copy (Make copy of table)
- Page-oriented NL-join, Block NL-join, Indexed NL-join, Hash join, Sort-Merge join
- Sort
- Bulk-load table from file
- Insert, delete, update
- Collect/update statistics
- Check some integrity constraint
- ...

27

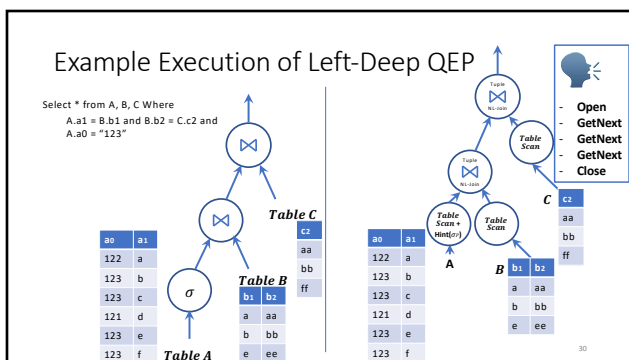
27



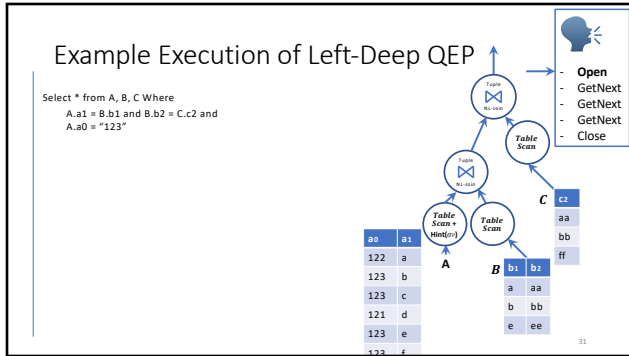
28



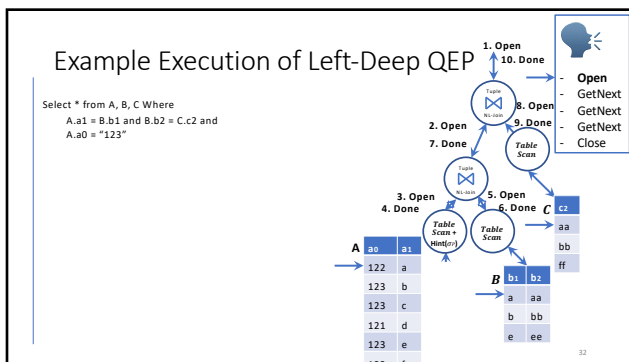
29



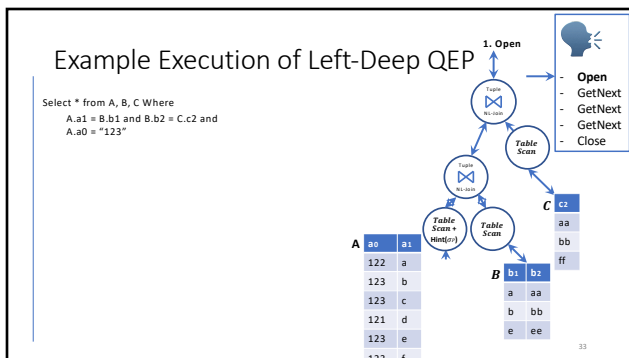
30



31



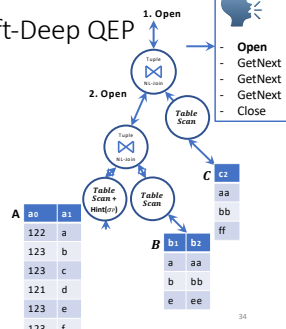
32



33

Example Execution of Left-Deep QEP

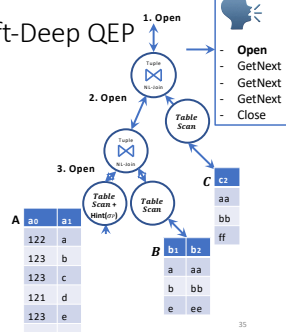
Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"



34

Example Execution of Left-Deep QEP

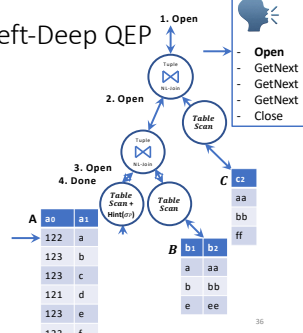
Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"



35

Example Execution of Left-Deep QEP

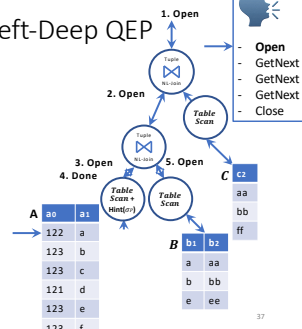
Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"



36

Example Execution of Left-Deep QEP

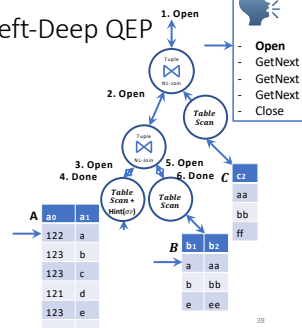
Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"



32

Example Execution of Left-Deep QEP

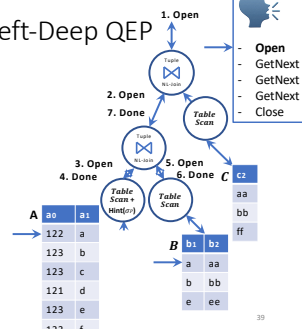
Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"



38

Example Execution of Left-Deep QEP

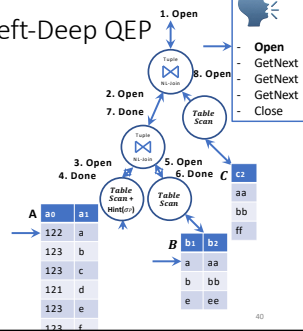
Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"



39

Example Execution of Left-Deep QEP

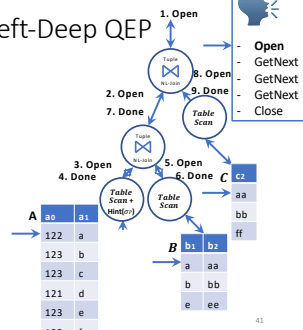
Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"



40

Example Execution of Left-Deep QEP

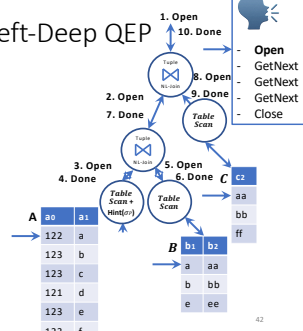
Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"



41

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"



42

Example Execution of Left-Deep QEP

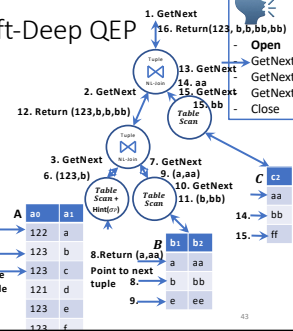
Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

9. (a,aa) does not qualify the
join predicate A.a1 = B.b1

14. aa does not qualify the join
predicate B.b2 = C.c2

4. Does not qualify
The Select Predicate
(Hint) A.a0 = "123"
122 ≠ 123

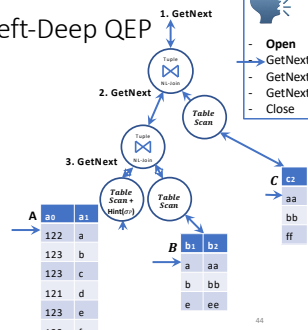
5. Advance to next tuple.
Qualifies Select Predicate
Return (123,b) and Advance
Cursor to point to next tuple



43

Example Execution of Left-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

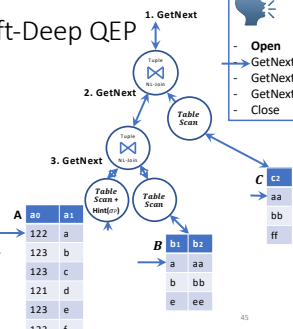


44

Example Execution of Left-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

4. Get Tuple (122,a)
Does not qualify
the Select Predicate
(Hint) A.a0 = "123"
122 ≠ 123
Advance to next tuple.

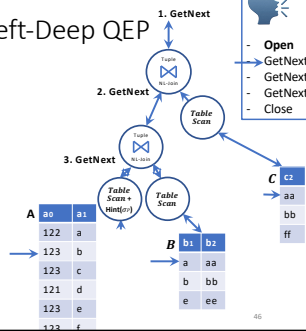


45

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 $A.a1 = B.b1$ and $B.b2 = C.c2$ and
 $A.a0 = "123"$

5. Get Tuple (123,b)
 It qualifies Select Predicate
 Return (123,b)

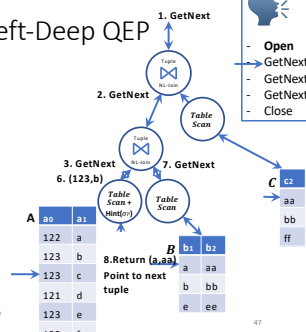


46

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 $A.a1 = B.b1$ and $B.b2 = C.c2$ and
 $A.a0 = "123"$

5. Get Tuple (123,b)
 It qualifies Select Predicate
 Return (123,b) and Advance
 Cursor to point to next tuple



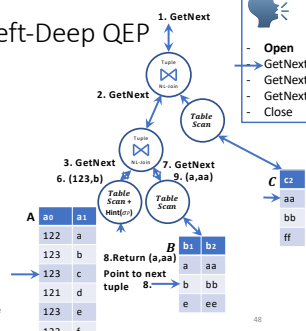
47

Example Execution of Left-Deep QEP

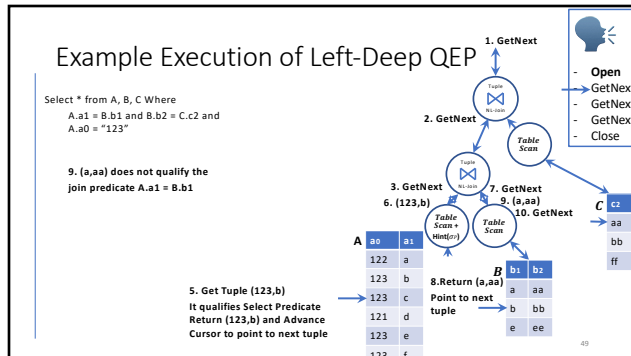
Select * from A, B, C Where
 $A.a1 = B.b1$ and $B.b2 = C.c2$ and
 $A.a0 = "123"$

9. (a,aa) does not qualify the
 join predicate $A.a1 = B.b1$

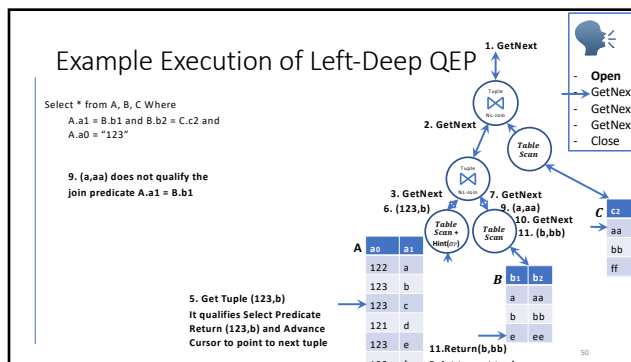
5. Get Tuple (123,b)
 It qualifies Select Predicate
 Return (123,b) and Advance
 Cursor to point to next tuple



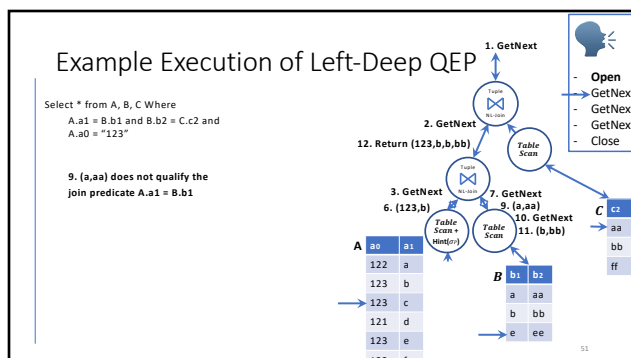
48



49



50

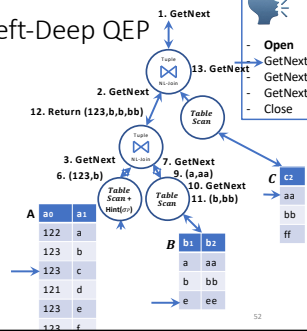


51

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

9. (a,aa) does not qualify the
 join predicate A.a1 = B.b1



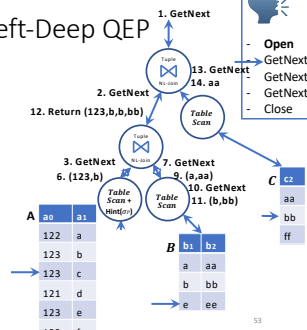
52

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

9. (a,aa) does not qualify the
 join predicate A.a1 = B.b1

14. aa does not qualify the join
 predicate B.b2 = C.c2



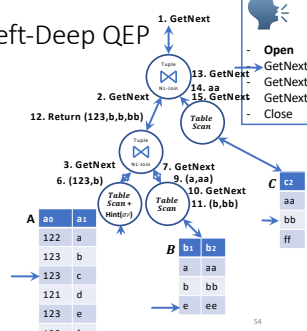
53

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

9. (a,aa) does not qualify the
 join predicate A.a1 = B.b1

14. aa does not qualify the join
 predicate B.b2 = C.c2



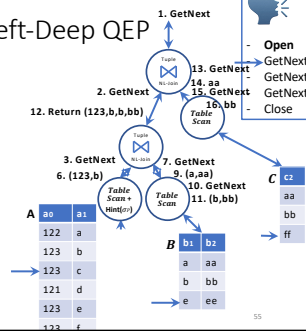
54

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

9. (a,aa) does not qualify the
 join predicate A.a1 = B.b1

14. aa does not qualify the join
 predicate B.b2 = C.c2



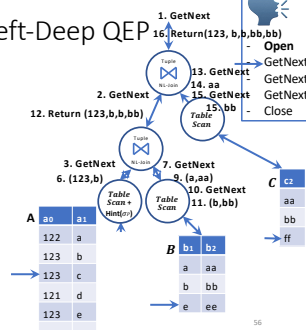
55

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

9. (a,aa) does not qualify the
 join predicate A.a1 = B.b1

14. aa does not qualify the join
 predicate B.b2 = C.c2

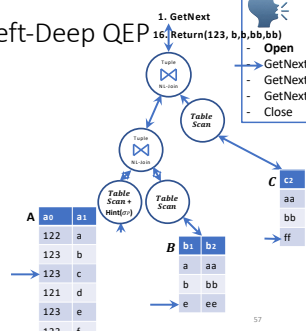


56

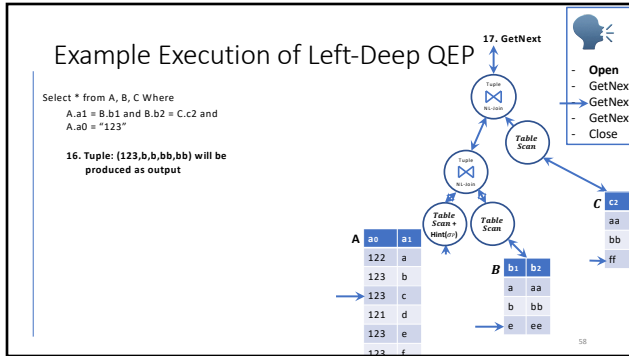
Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

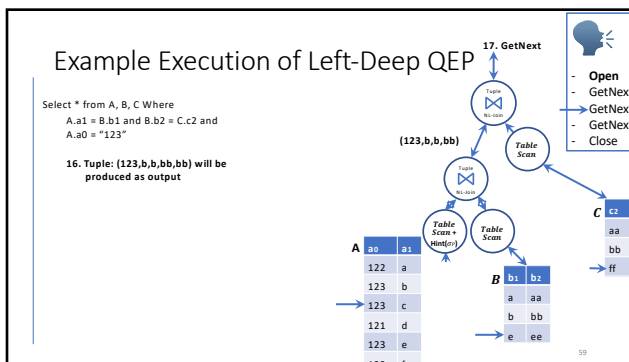
16. Tuple: (123,b,bb) will be
 produced as output



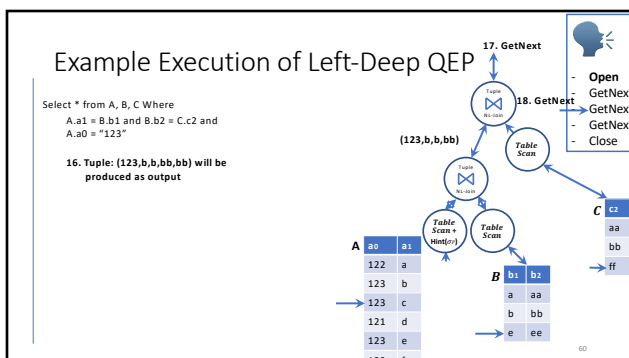
57



58



59



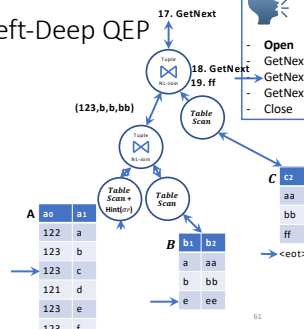
60

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate



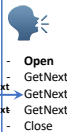
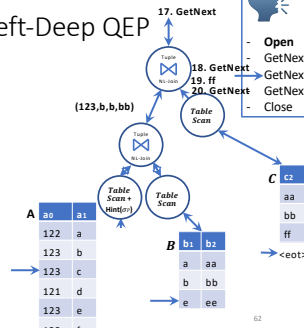
61

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate



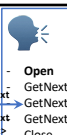
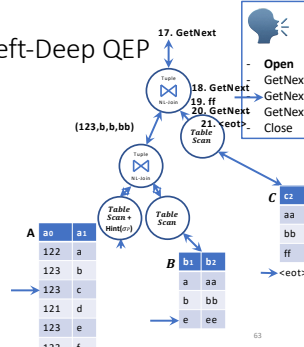
62

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate



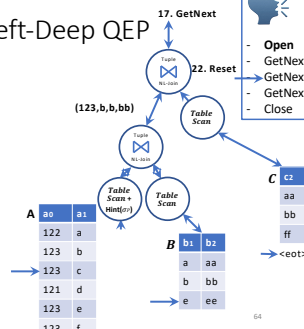
63

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate



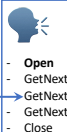
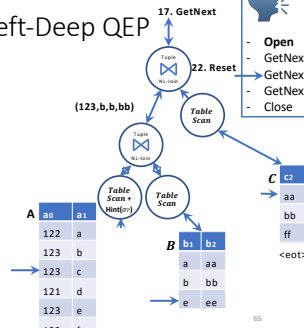
64

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate



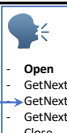
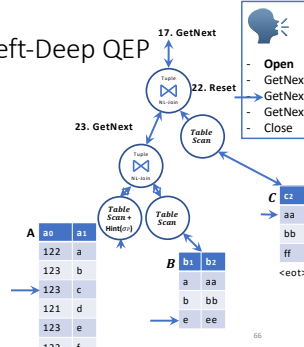
65

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate



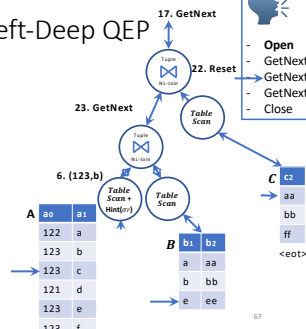
66

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate



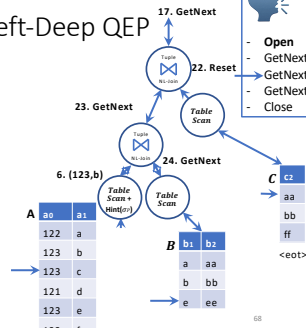
67

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate



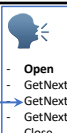
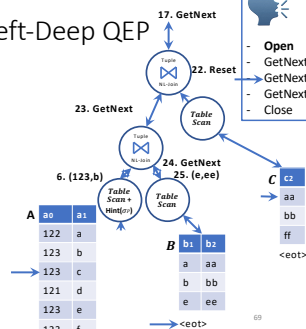
68

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

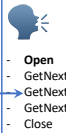
16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate



69

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



- Open
- GetNext
- GetNext
- GetNext
- Close

[illegible]

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



- Open
- GetNext
- GetNext
- GetNext
- Close

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



- Open
- GetNext
- GetNext
- GetNext
- Close

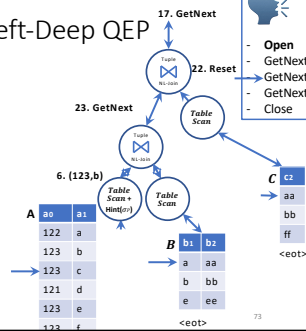
Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



Open
 - GetNext
 - GetNext
 - GetNext
 - Close

73

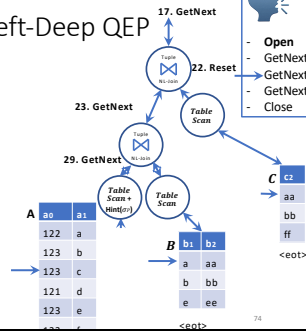
Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



Open
 - GetNext
 - GetNext
 - GetNext
 - Close

74

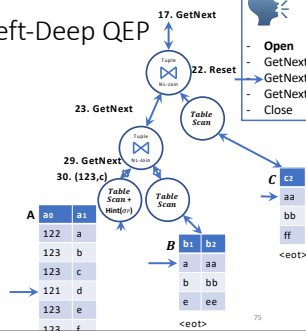
Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



Open
 - GetNext
 - GetNext
 - GetNext
 - Close

75

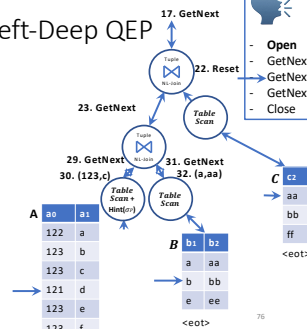
Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



76

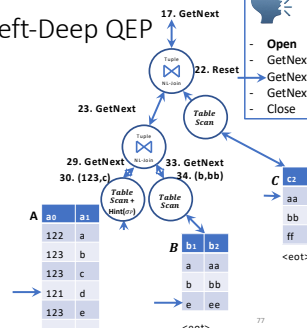
Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



77

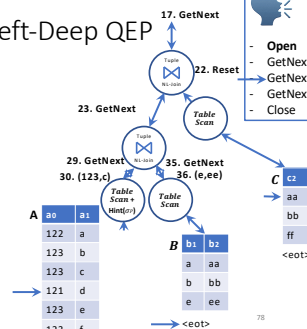
Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



78

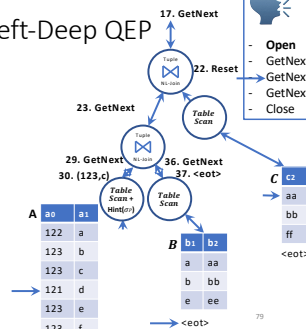
Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



79

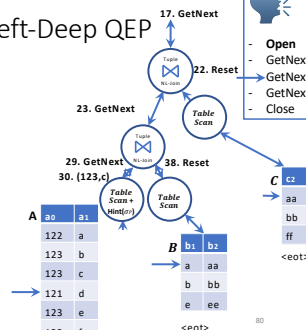
Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate



80

Example Execution of Left-Deep QEP

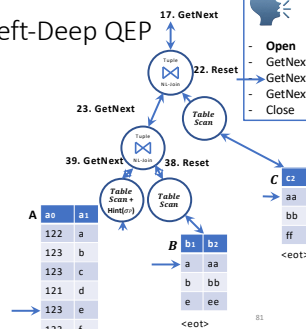
Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate

40. Tuple (121,d) does not satisfy the select predicate



81

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

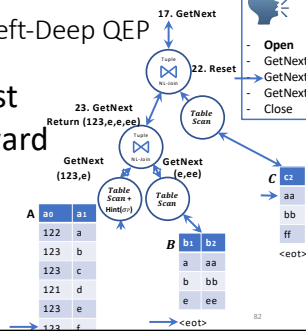
16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate

40. Tuple (121,d) does not satisfy the select predicate

Fast Forward



Open
 - GetNext
 - GetNext
 - GetNext
 - Close

82

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

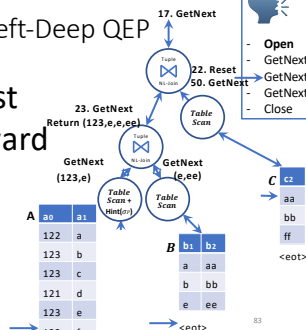
16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate

40. Tuple (121,d) does not satisfy the select predicate

Fast Forward



Open
 - GetNext
 - GetNext
 - GetNext
 - Close

83

Example Execution of Left-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

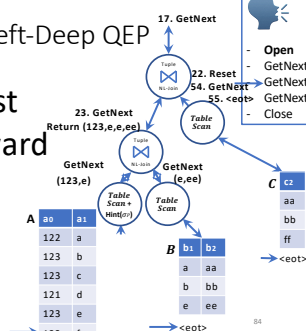
16. Tuple: (123,b,b,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate

40. Tuple (121,d) does not satisfy the select predicate

Fast Forward



Open
 - GetNext
 - GetNext
 - GetNext
 - Close

84

Example Execution of Left-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

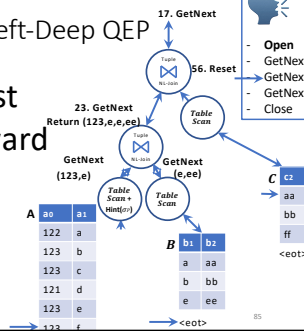
16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate

40. Tuple (121,d) does not satisfy the select predicate

Fast Forward



85

Example Execution of Left-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

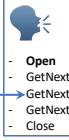
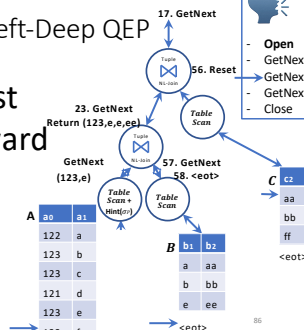
16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate

40. Tuple (121,d) does not satisfy the select predicate

Fast Forward



86

Example Execution of Left-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

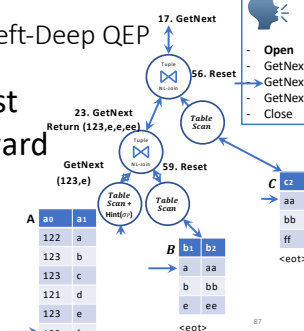
16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate

40. Tuple (121,d) does not satisfy the select predicate

Fast Forward



87

Example Execution of Left-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

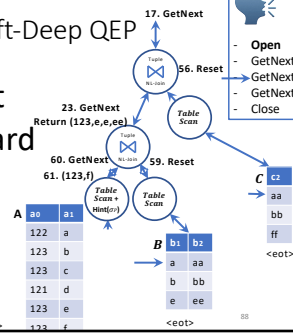
16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate

40. Tuple (121,d) does not satisfy the select predicate

Fast Forward



88

Example Execution of Left-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

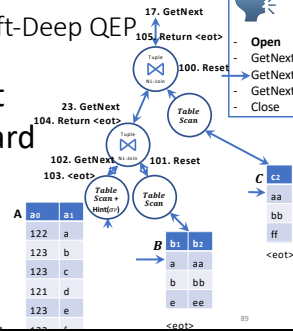
16. Tuple: (123,b,b,bb,bb) will be produced as output

19. Tuples (123,b,b,bb) and (ff) do not qualify the join predicate

25. Tuples (123,b) and (e,ee) do not qualify the join predicate

40. Tuple (121,d) does not satisfy the select predicate

Fast Forward

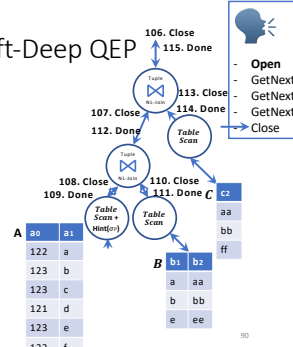


89

Example Execution of Left-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

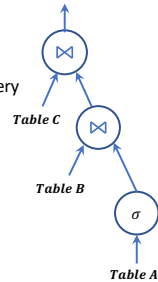
16. Tuple: (123,b,b,bb,bb) will be produced as output



90

Right-Deep QEPs

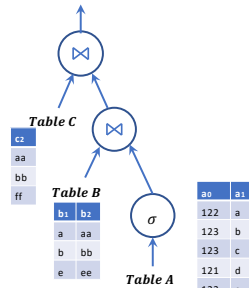
- Does a Right-deep QEP differ in terms of query processing w.r.t. Left-deep QEPs?



91

Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

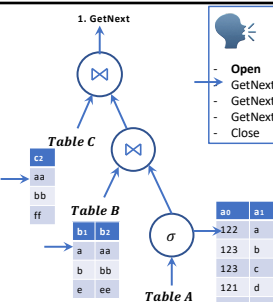


92

Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

- Assume we executed the Open method all the way
- Now, we execute GetNext



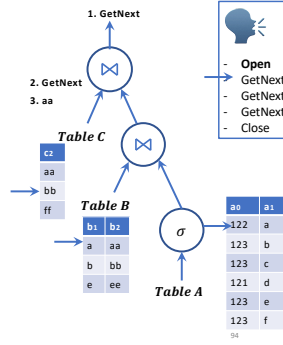
1. GetNext

- Open
- GetNext
- GetNext
- GetNext
- Close

93

Example Execution of Right-Deep QEP

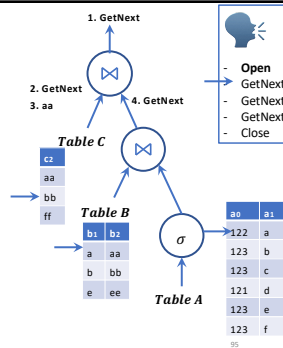
Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"



94

Example Execution of Right-Deep QEP

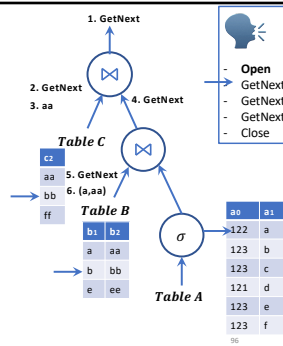
Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"



95

Example Execution of Right-Deep QEP

Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"

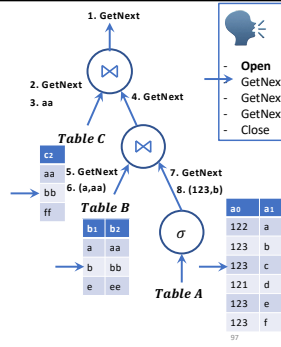


96

Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

Tuple (122,a) does not qualify the select predicate



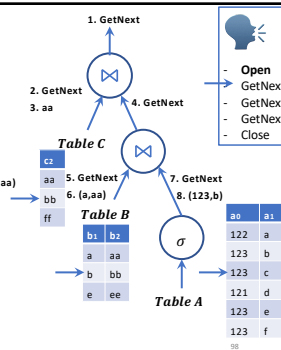
97

Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

7.1. Tuple (122,a) does not qualify the select predicate

9. All other tuples in Table A will not join with tuple (a,aa)



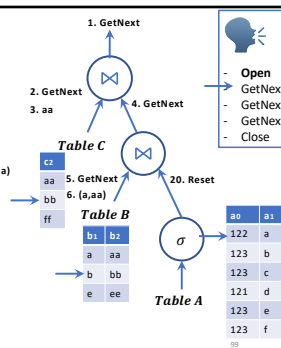
98

Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

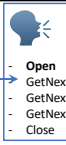
- All other tuples in Table A will not join with tuple (a,aa)

- Fast Forward



99

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"



Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"



Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"



54. Tuples (aa) and (e,ee,123,e) predicate

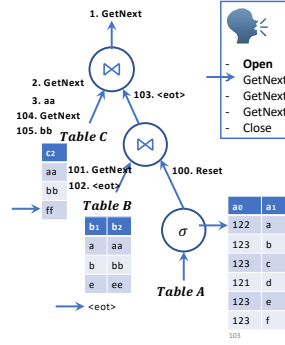
34

Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

106. Now, the join with Tables B and A will need to be reevaluated for every new tuple of C

107. Fast Forward



103

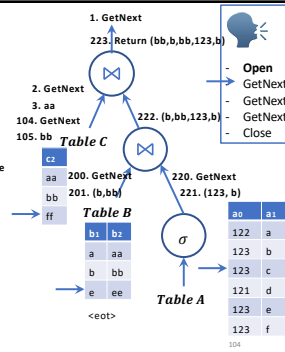
Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

222. Tuples (b,bb) and (123,b) satisfy the join predicate for the second time (Recall Step 32)

223. Tuples (bb) and (b,bb,123,b) satisfy the join predicate and is produced as output

224. Fast Forward



104

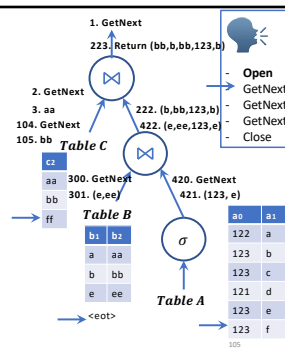
Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

423. Tuple (e,ee,123,e) got reevaluated for the second time (Recall Step 53)

424. Tuple (bb) and (e,ee,123,e) do not satisfy the join predicate

425. Fast Forward to the end as no more tuples will qualify

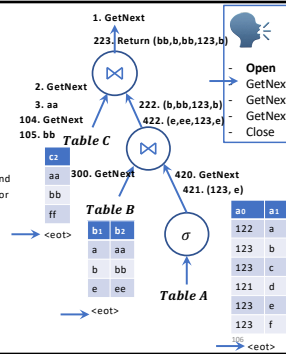


105

Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

Contrast the two snapshots of execution in this Slide and the snapshot in Slide 102 (listed also in the next slide for simplicity



106

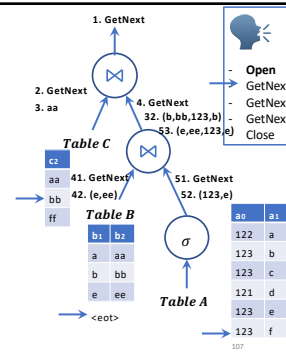
Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

Fast Forward

54. Tuples (aa) and (e,ee,123,e) do not satisfy the join predicate

55. Fast Forward



107

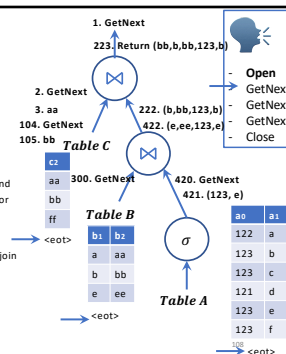
Example Execution of Right-Deep QEP

Select * from A, B, C Where
A.a1 = B.b1 and B.b2 = C.c2 and
A.a0 = "123"

Contrast the two snapshots of execution in this Slide and the snapshot in Slide 102 (listed also in the next slide for simplicity

Notice that the join of Tables A and B was performed twice, and of course in each time, produced the same join result: Tuples: (b,bb,123,b) and (e,ee,123,e)

Can you suggest an optimization here?

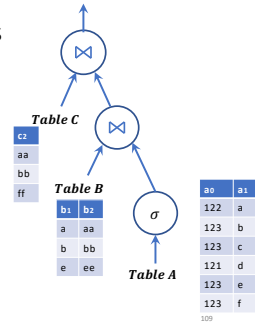


108

Optimizing Right-deep QEPs

- Will add a temp operator on top of each join operator:
 - Perform the once and save its results
 - Next time, instead of re-evaluating it, access the temp table

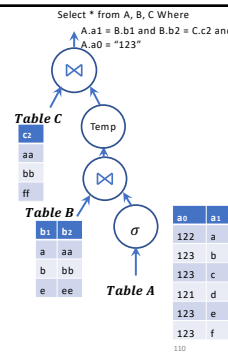
Select * from A, B, C Where
 A.a1 = B.b1 and B.b2 = C.c2 and
 A.a0 = "123"



109

Optimizing Right-deep QEPs

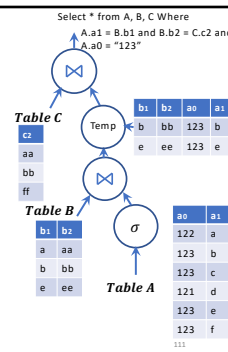
- Note the addition of the Temp operator above the join of B and A
- Temp will create a temporary table that is the result of the subtree below it
 - Temp = B ⋈ B.b1 = A.a1 (σ_{A.a0 = "123"}(A))
- The temporary table is created the first time the Temp operator is invoked



110

Optimizing Right-deep QEPs

- Note the addition of the Temp operator above the join of B and A
- Temp will create a temporary table that is the result of the subtree below it
 - Temp = B ⋈ B.b1 = A.a1 (σ_{A.a0 = "123"}(A))
- The temporary table is created the first time the Temp operator is invoked when processing Tuple (aa) of Table C
- For the remaining tuples in C, the temporary table is accessed



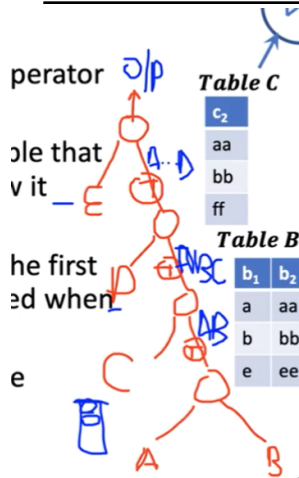
111

How to optimize right deep QEP

By adding a temporary table (result of subtree below it) above a join operator. So we don't need to iterate to get every join permutation.

i.e. If there is a query that E joins D joins C joins A and B.

We add a temporary table above each join operator



Comparison between Left- and Right-deep QEPs

- Left-deep Trees are the more favorable
- Allow for non-blocking operators
- Do not need to write temporary tables on disk
- Do not have to read the temporary tables from disk over and over for every tuple of the outer tables
- Make the query optimization job easier as we do not have to search over all possible trees that can evaluate the query
- But still exponential as we need to know the order of the joins (which join to perform first).

112

However, temporary table is blocking (get all results from subtree until done) and need space to store temporary results
